

Introduction to Version Control and Git

Dave McEwan
Tomorrow 99th May 2099

Introduction

Agenda

-1100 Meet and greet.
- 1100..1115 Introduction
- 1115..1200 Background: Revision control, version numbering.
- 1200..1300 Practical/interactive: [RCS](#) and [git](#) mechanics.
- 1300..1330 Lunch.
- 1330..1430 Local setup: GitLab, SemVer, Sub-modules, CI, etc.
- 1430..1600 Practical/interactive: Merging.

Assumptions

- You are comfortable working with Bash (or Zsh, Tcsh) on Linux.
- You have experience in collaboratively developing code.
- You have experience with some sort of version control before.

... and that you have setup these things before today:

- WiFi access.
- `/home/<username>/` with `rwX` permissions.
- Local or SSH access to Linux machine.

Objectives

- Understand fundamental mechanics of git.
- Be comfortable exploring the internals of git repositories.
- Know how to dig yourself out of common problems with git.
- Understand interactions between our systems.

Roundtable Introductions

- Who are you?
- Who am I?

Background

diff and patch

- Basis for all version control systems discussed today.
- POSIX since X/Open Portability Guide Issue 4 (1992).
- `diff` takes file1+file2, reports their differences in a patchfile.
- A patchfile precisely describes differences between two files.
 - Most common modern format is “unified” (vs “ed”, “context”).
 - Some context is often required.
 - Differences are identified line-by-line.
- `patch` takes a file+patchfile, applies the differences in-place.
- Inputs must be vaguely similar.

Garbage in, garbage out.

RCS Overview

- <https://www.gnu.org/software/rcs/>
- 1st generation, but not obsolete!
- Current version is 5.10.1 (2022-02-02).
- Implemented as a collection of utilities (not a single exe):
 - `ci, co`
 - `ident, rcs, rcsclean, rcsdiff, rcsfreeze, rcsmerge, rlog`

RCS Important Details

- Stored as separate/unmerged reverse deltas.
 - I.e. delta says how to get back to *previous* version.
 - Predecessor SCCS used forward deltas saying how to get to *next* version.
- Lock indicates that somebody *intends* to deposit a newer revision.
- Branches are on version numbers
 - Symbolic name (similar to a “tag”) is a prefix shortcut.
- Concept of joining is similar to merging.
- Stamping, like `$Header$`, expands strings on checkin.
- An edit script is a patchfile.

CVS Overview

- <https://cvs.nongnu.org/>
- 2nd generation, mostly obsolete.
- Frontend to RCS which adds a client/server model.
- Centralised repository, usually accessed over RSH (predecessor to SSH).
- Composed of fewer executables (`cvs`, `cvspserver`).
- Delta compression distinguishes between text and binary.
- A “project” is a set of related files.
- Introduced “loginfo” similar to git’s hooks.

SVN Overview

- <https://subversion.apache.org/>
- 2nd generation, mostly obsolete.
- Intended to be successor to CVS, and mostly compatible.
- Many implementations of both clients and servers.
- Centralised repository, accessed over SSH, HTTP, or HTTPS
- Atomic commits to multiple files.
- Properties (key/value pairs) expand the concept of stamping.
- The SVN tool itself is version controlled
 - There are incompatibilities between versions (and that's ok).
- Relies heavily on conventions (trunk, branches, tags) so it's easy to make a mess.
- GUIs are all 3rd-party.

Git Overview

- <https://git-scm.com/>
- 3rd/current generation.
- Implemented as a single executable (`git`).
- Distributed network of repositories, accessed over SSH or HTTPS.
- Designed to support development of Linux kernel.
 - Thousands of developers, distributed globally, with restrictions on network access, on all sorts of operating systems.
- Every repository has a full copy of history.
- Concepts of branches and tags are first-class citizens with precise definitions.
- GUIs provided in standard distribution, but many 3rd-party tools exist.
- Locking concept doesn't translate well.
 - Close, but slightly different, approximations can be emulated via hooks, branches, and talking to each other.

Git Commits

- Each commit has a unique SHA1 hash.
- Concept of version numbering is intentionally avoided.
- Concept of stamping, like `$Header$`, doesn't exist.
 - Behaviour can be emulated via tags and hooks.
- Tags can be any format you like (please use SemVer).
 - “lightweight” tag is a convenient alias for a commit hash.
 - “annotated” tag is an object in its own right with a checksum, tag message, tagger's name, date, GPG signature.
 - In the same way as branches, tags must be pushed individually.
- Branches are fast, much more convenient than SVN and its predecessors.

Git Hooks

- Run scripts on your local repo on specific events.
- Local events: `pre-commit`, `commit-msg`, `post-merge`, `pre-push`, and more ...
- Remote/server events: `pre-receive`, `update`, `post-receive`
- Must be setup per repository.

SemVer and SemVerSoC

- <https://semver.org/>
- `<major>.<minor>.<patch>-<pre-release>+<build>`
- Version numbers and the way they change have convey well-defined meaning - that tools can rely on.
- Used for many/most open-source projects.
- ... See the tiny webpage. ...
- Releases may use [SemVerSoC](#).

- CI: Continuous Integration
 - Run regression on specific events, and complain (loudly, via email) if something is broken.
 - “Push branch X”, “Commit to branch Y”, “Every Monday at 1200”
 - GitHub has [Actions](#).
 - Great for anything code-based.
- CD: Continuous Deployment
 - If the new version passes regression, then deploy it.
 - Can be built into CI flow.
 - GitHub has [Dependabot](#).
 - Great for web development, but not for hardware or critical software.

Practical/Interactive: RCS and Git Mechanics

Preliminary

1. Let's create a space for all of today's practical work.

- `cd /work/${USER}/`
- `mkdir -p VCG/morningRCS`
- `mkdir -p VCG/morningGit VCG/morningSomewhereElse`
- `mkdir -p VCG/afternoon`

2. `unset noclobber` or `set +o noclobber`

3. Look inside `~/.gitconfig` if it exists.

RCS 1of4

- Objective 1: Feel familiar with inspecting the system.
- Objective 2: Give a point of comparison against git.
- Estimated time: 10 minutes.
- There's a similar tutorial here:
<https://www.madboa.com/geek/rcs/>

RCS 2of4

Let's start by making a couple of files to play with.

- `cd VCG/morningRCS/`
- `printf 'hello\nworld\n' > foo.txt`
- `printf 'red\nblue\n' > bar.txt`
- `ls -al ./ RCS/`

And provide a place for RCS to keep its data.

- `mkdir RCS`

RCS 3of4

Next, let's tell RCS to manage these files. By default, RCS will delete the original files on checkin, so you usually want the `-u` option to keep them in the working directory.

- `ci -u foo.txt bar.txt`
- `ls -al ./ RCS/`
- `cat RCS/foo.txt,v`

RCS 4of4

Now, let's try making a change to `foo.txt`

- `echo writableNo >> foo.txt`

But wait! You need to checkout first.

- `co -l foo.txt`
- `echo writableYes >> foo.txt`
- `ls -al ./ RCS/`

And finally, we can observe that we've made changes.

- `rcsdiff`

Git 1of11

- Objective 1: Feel familiar with inspecting the system.
- Objective 2: Be comfortable starting a new repository.
- Estimated time: 30 minutes.
- Follows then extends what we just did with RCS.

Git 2of11

Let's start by making a couple of files to play with...

- `cd VCG/morningGit/`
- `printf 'hello\nworld\n' > foo.txt`
- `printf 'red\nblue\n' > bar.txt`
- `ls -al`

... and initialising the repository.

- `git init`

Git 3of11

Let's have a look what that did to our workspace.

- `find . | sort`

Immediately noticable is that the management structure looks more like a database. Git commands mostly work on these files, so let's see an example of that with `.git/config`.

- `cat .git/config and/or git config user.name`
- `git config user.name "Sam Smith"`
- `cat .git/config and/or git config user.name`

Git 4of11

Next, let's tell git to manage our files. This has two steps (add to the staging area, then the atomic commit).

- `git add foo.txt bar.txt`
- `git commit -m 'Little message about the changes.'`

Also different from RCS is that we can always edit our working files.

- `echo writableYes >> foo.txt`

Git 5of11

We can see changes between the working directory and the staging area (also known as the cache).

- `git diff`

And changes between the staging area and the committed files.

- `git add foo.txt`
- `git diff --cached`

Git 6of11

The actions of tracking files and viewing differences between revisions in Git are immediately comparable to RCS.

- Can see exactly what is meant by atomicity?
- Can you see how things are arranged in your local workspace?

Now let's have a look at the distributed and decentralised features. Create a *bare* repository that can work like GitLab/BitBucket/GitHub, then have a look inside. Do this somewhere else on the filesystem - not in the repo you've just created.

- `cd VCG/morningSomewhereElse/`
- `git init --bare MyRepo`
- `ls MyRepo/`

Git 8of11

Back in our original repo, we can add that as a remote for our original. We're using the name `myremote`, but BitBucket and GitHub suggest the name `origin` in their documentation. You can use whatever alias you like - it's your choice and doesn't affect anybody else.

- `cd VCG/morningGit/`
- `git remote add myremote path/to/MyRepo`

Note that this does nothing to MyRepo yet. When we push our original repo, that's when the network access occurs.

- `git push myremote master`

Now have a look around MyRepo where you can see that the contents of `objects/` is identical to those in our local copy. You can add as many remotes as you like, name them however you like, and synchronise them however you like.

To see a nice overview, you can use the standard GUIs for working with local changes and seeing what your local repo knows about other repos. Let's have a quick look now. Both `git gui` and `gitk`, implemented in Tcl/Tk, are included in the usual distributions of Git and provide a consistent interface on Linux, BSDs, MacOS, and Windows.

- `git gui &`
- `gitk --all &`

Git 10of11

Finally, let's play with a merge which has a little conflict. Merging the changes from a branch in your local repo is the same process as merging the changes from a branch on a remote repo - just fetch first.

- `git checkout -b firstBranch`
- `echo apples >> foo.txt && git add . && git commit -m 'Add apples.'`
- `git checkout master && git checkout -b secondBranch`
- `echo oranges >> foo.txt && git add . && git commit -m 'Add oranges.'`

Finally, return our working directory to the HEAD of `master`.

- `git checkout master`

We'll try merging the first branch, then the second.

- `git merge firstBranch`
- `git merge secondBranch`

That didn't work because, in the second merge, changes had been made to the same part of the file touched by a commit since their common ancestor. We could also have made this conflict by adding the “apples” line directly on the `master` branch.

To fix the conflict, we need to choose what is right. The simplest way to do this is to accept either change using `git gui`, then modify it manually before the merge is committed. We'll go over different ways of merging in the afternoon.

Local Setup

Outline

- GitLab, BitBucket, GitHub
- SemVer and SemVerSoC
- Jenkins
- Q and A

GitLab, BitBucket, GitHub

- Hooks may prevent rewriting history, like amending pushed commits.
- Hooks may enforce branch our naming scheme.
- Markdown is rendered on all web views.
 - README, commit messages, pull requests
- PDF is also viewable in the web browser, but not DOCX.
- SVG,PNG,JPG are rendered, but not VSDX.
 - SVG can be diffed – [Inkscape](#) is recommended for most diagrams and [Wavedrom](#) for waveforms.
 - Visio obfuscates SVG!

Git Submodules

- Incorporate other git repositories into your repository.

SemVer and SemVerSoC

- **SemVer** ascribes meaning to software version numbers.
- **SemVerSoC** clarifies SemVer's rules for HDL/RTL projects.
- You've probably used SemVer constraint solvers before - in package managers like `apt` (Debian), `rpm` (RHEL/CentOS), `cargo` (Rust), `pip` (Python).

Jenkins

- <https://github.com/jenkinsci>
- A continuous integration tool.
- Free, open-source software written in Java.

Practical/Interactive: Git Merging

Outline

- Preliminaries
- Example 1: Diffing with tkdiff.
- Example 2: Merging without change.
- Example 3: Merging without conflict.
- Example 4: Merging with conflict.
- Example 5: Merging with kdiff3.

Preliminaries

1. Examples with FSF's `reuse-tool`.
2. Non-merge vs merge commits.
3. GUIs.
4. Stashing.
5. `reset` vs `revert`.
6. Commit message format.
7. `fetch`, `merge`, `rebase`, and `pull`.
8. Nordic's BitBucket branch naming scheme.

Preliminary 1of7: Examples with FSF's reuse-tool

- Create a new directory to work on the repository.
 - `cd VCG/afternoon/`
 - `git clone https://github.com/fsfe/reuse-tool.git`
 - `cd reuse-tool`
- Now we can see some history.
 - `git log`
- As well as the branches and remotes.
 - `git branch`
 - `git remote -v`
 - `git branch -a`

Preliminary 2of7: Non-Merge vs Merge Commits

- There two types of commits in Git: non-merge, and merge.
- Non-merge commits are the usual kind.
 - Extend a branch's history.
 - Single ancestor.
 - Must contain differences.
 - Created with `git commit`.
- Merge commits join two branch's histories.
 - Two ancestors.
 - *May* contain differences.
 - Created with `git merge`.
- Have a look at merge commits: `git log`
 - A merge commit with no differences: `git show 95b6dd09`
 - A merge commit with differences: `git show 65b3206a`

Preliminary 3of7: GUIs

- CLI will give you canonical status/results, but GUIs can make some tasks easier.
- `git status`: Visualise status with `git gui`.
 - In the default distribution.
 - Also useful for amending the previous commit.
- `git log`: Visualise branch history with `gitk --all`.
 - In the default distribution.
- `git diff`: Navigate diffs one file at a time with `git difftool`.
 - `git config --global diff.tool tkdiff`
- `git merge`: Resolve conflicts with `git mergetool`.
 - `git config --global merge.tool kdiff3`

Note: `--global` changes `~/.gitconfig`, instead of `./.git/config`

Preliminary 4of7: Stashing

- A handy tool for digging yourself out of a mess.
- Save uncommitted and unstaged changes temporarily.
- Easier than saving files to temporary locations.
- Implemented as a stack.
 - Push unstaged changes with `git stash`
 - Pop with `git stash pop`
- Let's try that quickly by making a change, then attempting to pull.
 - `echo BREAKTHINGS >> ./README.md`
 - `git pull --rebase`
- For a rebasing pull to succeed, tracked files must be unmodified. Save our change onto the stash stack.
 - `git stash`
 - `git pull --rebase`
- Now we can get our change back.
 - `git stash pop`

Preliminary 5of7: `reset` vs `revert`

- To undo unwanted changes that aren't committed, reset all tracked files.
 - `git reset --hard`
- That puts your working repo back to `HEAD`.
- To undo changes made by existing commits, we need a new commit which reverses the diff.
 - `git revert <commit>`
- To help yourself and colleagues:
 - Write useful commit messages in the de-facto standard form.
 - Use searchable keywords in commit messages.
 - Review the diff before you make each commit.
 - Keep unrelated changes in separate commits.

Preliminary 6of7: Commit Message Format

- De-facto format supported by GitLab, BitBucket, GitHub, 3rd-party tools, etc.
- First line is a subject, up to 50 characters.
- Second line is empty/blank.
- Subsequent lines are the body, up to 72 characters.
- Subject is in imperative mood, i.e. commanding.
 - Good: “*Add* support for Foo.”
 - Bad: “*Adds* support for Foo.”
- Alternatively, if the commit is very simple and doesn't need a body, word the subject as a phrasal noun.
 - “Foo support”
- Both subject and body should be formatted in [Markdown](#).

Preliminary 7of7: Fetch, Merge, Rebase, and Pull

- `git fetch`: Fetch/copy the newest branch and commit data from a remote, but don't change any branches or working files.
- `git merge`: Performs a 3-way merge between two branches and their most recent common ancestor.
- `git rebase`: Take the series of commit patches and apply them to rewrite the branch's history.
 - Rewriting history is not allowed on some server setups. In those systems, you'll rarely (if ever) use the rebase command.
- `git pull`: Execute `git fetch` then `git merge` (or `git rebase`).
 - If you're unsure, use separate steps for fetching and merging.

Diffing with `tkdiff` 1of1

- Let's introduce a harmless change to a tracked file.
 - `echo '## HARMLESS' >> src/reuse/lint.py`
 - `git difftool`
- We can also use `tkdiff` to diff over branches.
 - `F=src/reuse/lint.py`
 - `B1=origin/master`
 - `B2=origin/summary-help`
 - `git difftool $B1:$F $B2:$F`
- First argument is LHS, second argument is RHS.
- Use `n` and `p` to move quickly to next and previous changes.
- See the documentation for all the ways you can diff.
 - <https://git-scm.com/docs/git-diff>
 - <https://git-scm.com/docs/git-difftool>

Merging Without Changes 1of3

- This is the simplest kind of merge, where the same changes have been made on converging branches.
- First, let's get back to a known state and make a change on one branch.
 - `git checkout master`
 - `git reset --hard`
 - `git checkout -b myBranchA`
 - `echo '## foo' >> README.md`
 - `git add !$`
 - `git commit -m 'Add foo comment, by Alice.'`

Merging Without Changes 2of3

- Second, get back to a known state and make the same change on another branch.
 - `git checkout master`
 - `git reset --hard`
 - `git checkout -b myBranchB`
 - `echo '## foo' >> README.md`
 - `git add !$`
 - `git commit -m 'Add foo comment, by Bob.'`

Merging Without Changes 3of3

- Now that our changes are committed, we can freely switch branches without losing work.
 - `git checkout myBranchA`
 - `git checkout master`
- To start the merge, first switch to the “destination” branch, i.e. the one that will still be worked upon. Let's choose to merge `myBranchA` into `myBranchB`.
 - `git checkout myBranchB`
 - `git status`
 - `git merge myBranchA`
 - Save and quit the editor
- Now, have a look at the history with `git log` and/or `gitk`.

Merging Without Conflict(s) 1of3

- This is also a simple kind of merge, where the completely different changes have been made on converging branches.
- First, let's get back to a known state and make a change on one branch.
 - `git checkout master`
 - `git reset --hard`
 - `git checkout -b myBranchC`
 - `echo '## foo' >> README.md`
 - `git add !$`
 - `git commit -m 'Add foo comment, by Alice.'`

Merging Without Conflict(s) 2of3

- Second, get back to a known state and make a different change on another branch.
 - `git checkout master`
 - `git reset --hard`
 - `git checkout -b myBranchD`
 - `echo '## foo' >> Makefile`
 - `git add !$`
 - `git commit -m 'Add foo comment, by Bob.'`

Merging Without Conflict(s) 3of3

- To start the merge, first switch to the “destination” branch, i.e. the one that will still be worked upon. Let's choose to merge `myBranchD` into `myBranchC`.
 - `git checkout myBranchC`
 - `git status`
 - `git merge myBranchD`
 - Save and quit the editor
- Now, have a look at the history with `git log` and/or `gitk`.

Merging With Conflict(s) 1of4

- This is the usual kind of merge, where the same changes have been made on converging branches.
- First, let's get back to a known state and make a change on one branch.
 - `git checkout master`
 - `git reset --hard`
 - `git checkout -b myBranchE`
 - `echo '## foo' >> README.md`
 - `git add !`
 - `git commit -m 'Add foo comment, by Alice.'`

Merging With Conflict(s) 2of4

- Second, get back to a known state and make a conflicting change on another branch.
 - `git checkout master`
 - `git reset --hard`
 - `git checkout -b myBranchF`
 - `echo '## bar' >> README.md`
 - `git add !$`
 - `git commit -m 'Add foo comment, by Bob.'`

Merging With Conflict(s) 3of4

- To start the merge, first switch to the “destination” branch, i.e. the one that will still be worked upon. Let's choose to merge `myBranchE` into `myBranchF`.
 - `git checkout myBranchF`
 - `git status`
 - `git merge myBranchE`
 - ... This needs some attention.

Merging With Conflict(s) 4of4

- Look at the conflicts with `git gui`.
- Search for conflicts with <<<<<<, =====, and >>>>>>.
 - That's 7 characters for each kind of conflict marker.
- The format is always the same.
 - Conflicts are marked beginning with <<<<<< (opening chevrons). and ending with >>>>>> (closing chevrons).
 - Each conflict has two parts, partitioned by ===== (the separator).
 - Lines outside of chevrons belongs to the most recent common ancestor (`master`).
 - First, lines between the opening chevrons and the separator are from the current/destination branch (`myBranchF`)
 - Second, lines between the separator and the closing chevrons are from the branch we want to merge in (`myBranchE`).
- Edit the file, commit, then look at the history.
- A simple strategy is to accept the “remote” version, then make edits, stage, then commit.

Merge Using `kdiff3` 1of3

- Let's repeat the previous example, but use a GUI tool for the merge step.
 - `git checkout master`
 - `git reset --hard`
 - `git checkout -b myBranchG`
 - `echo '## foo' >> README.md`
 - `git add !`
 - `git commit -m 'Add foo comment, by Alice.'`
 - `git checkout master`
 - `git checkout -b myBranchH`
 - `echo '## bar' >> README.md`
 - `git add !`
 - `git commit -m 'Add foo comment, by Bob.'`

Merge Using `kdiff3` 2of3

- To start the merge, first switch to the “destination” branch, i.e. the one that will still be worked upon.
 - `git checkout myBranchG`
 - `git status`
 - `git merge myBranchH`
 - `git mergetool`

Merge Using `kdiff3` 3of3

- There are 4 panes.
 - Left (Base) is the most recent common ancestor.
 - Middle (Local) is the current/destination branch (`myBranchG`)
 - Right (Remote) is the branch we want to merge in (`myBranchH`)
 - Bottom is the output of the merge.
- Controls and keyboard shortcuts are well described and intuitive.
- Right-click on conflicts in the bottom pane to select between solutions.
- You can type in the bottom pane.
- Save the output, then use `git gui` to view your changes as usual.

Q and A